Module 4A
Supervised Machine Learning

Different flavors of REGRESSION and General Linear Models

Review: Correlation:

- Measures the amount/degree of linear association between two **numerical** variables

- Estimate the degree to which variables **covary**
  - With no attempt to interpret the causality of the association

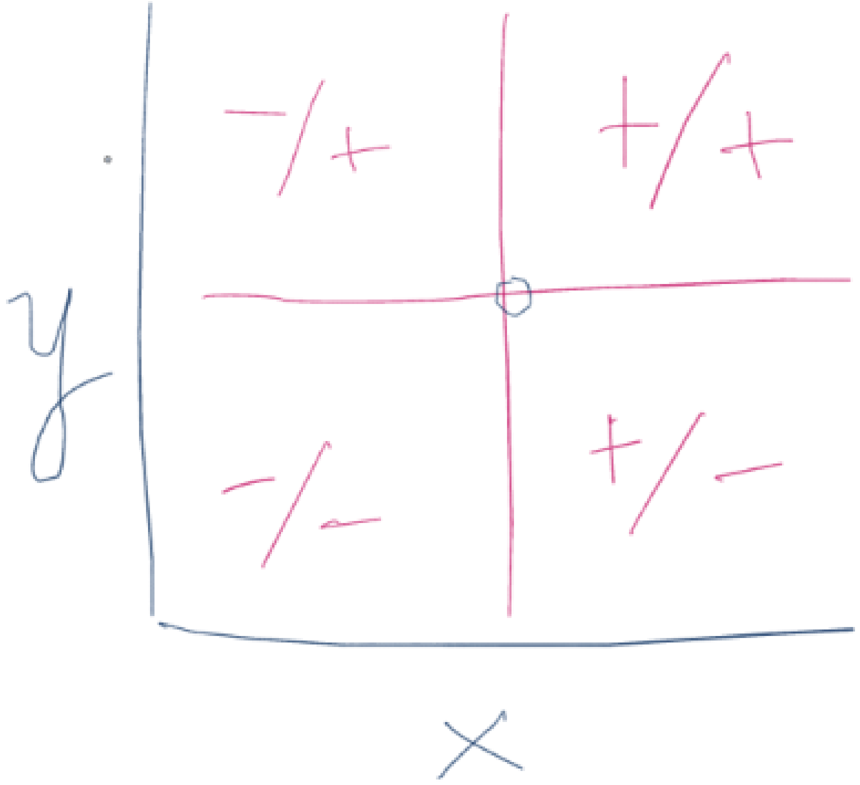

Example: arm length and leg length covary together (individuals with longer arms often have longer legs) but they are influenced by other underlying variables **not** each other (longer legs do not cause longer arms)

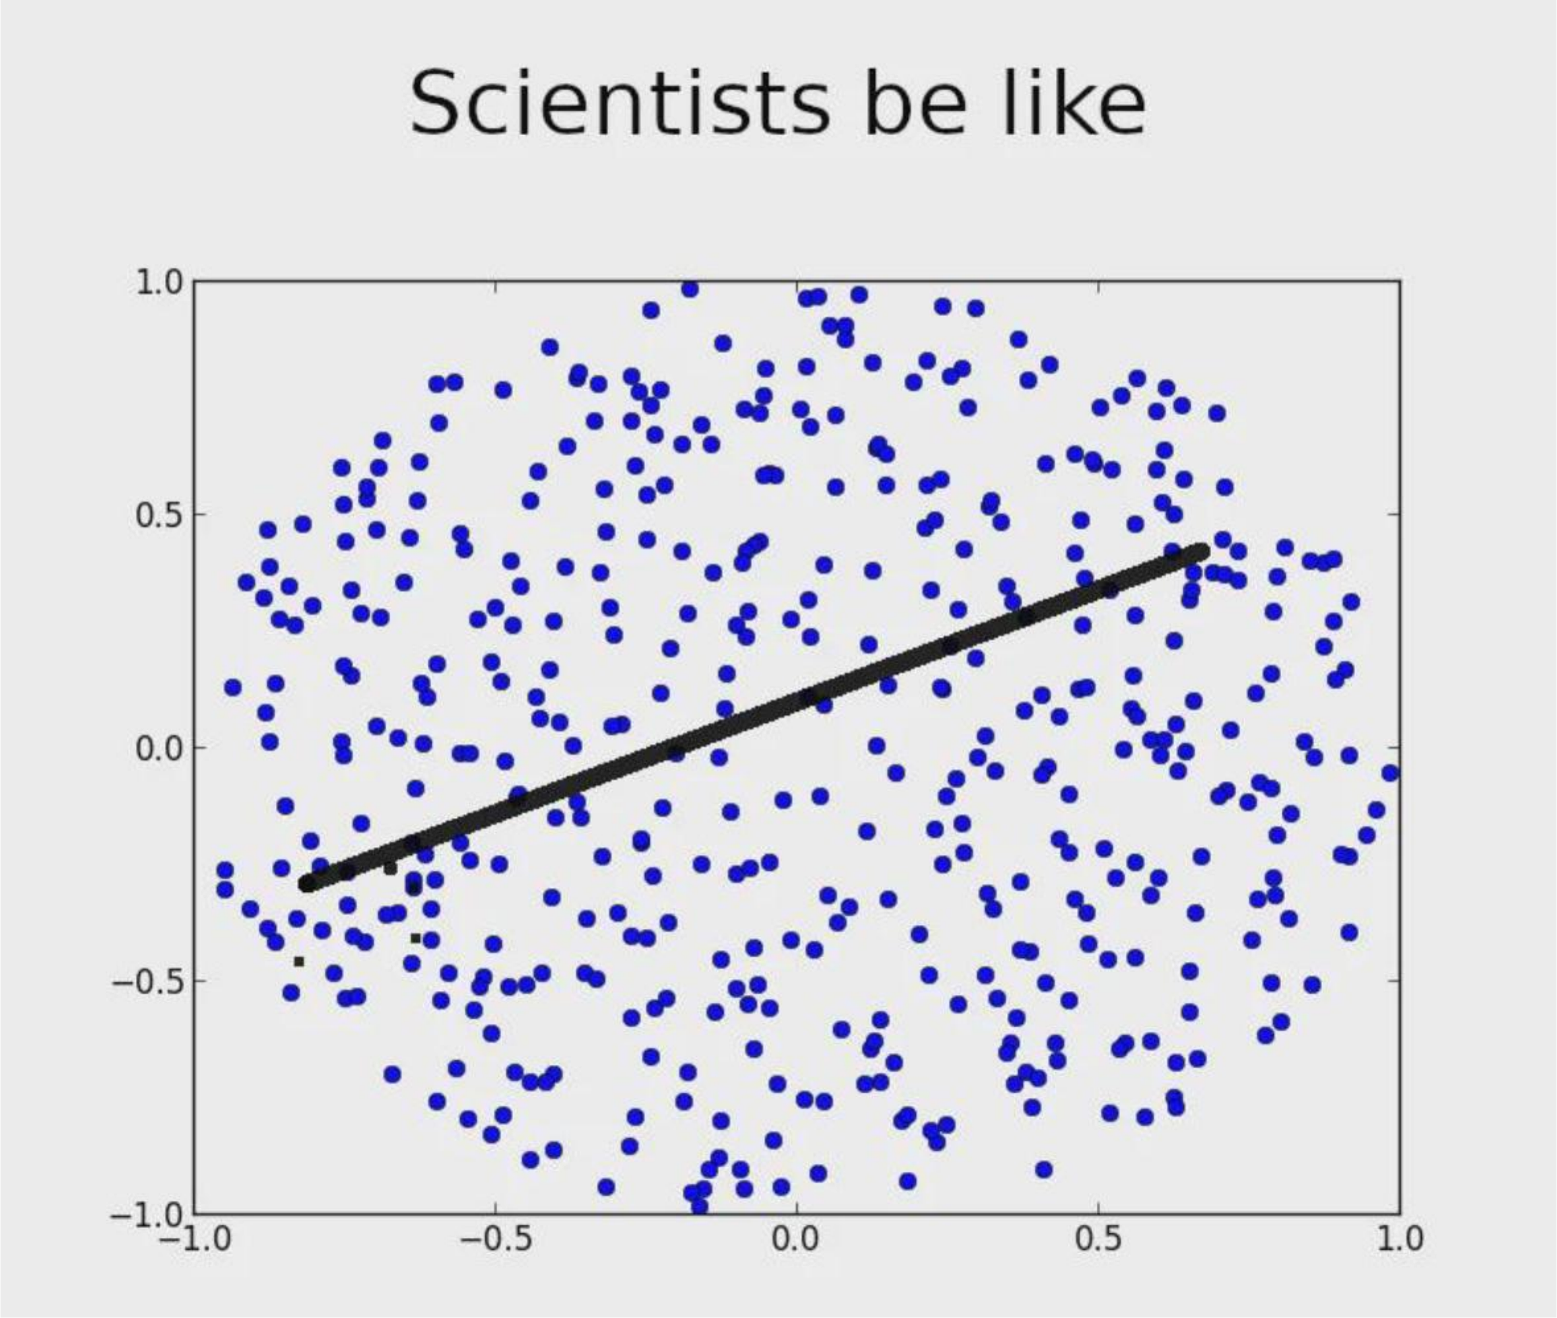
Scientists be like
1.0
0.5
0.0
−0.5
−1.0
−1.0
−0.5
0.0
0.5
1.0

Regression:

- Statistics is about prediction

- Used to **predict** value of one numerical variable from the value of another
  - predicting dependent/response variable, Y from independent/predictor X

- **Linear** regression assumes that the relationship between X and Y can be described by a line
  - Fits a straight line to a (messy) scatterplot

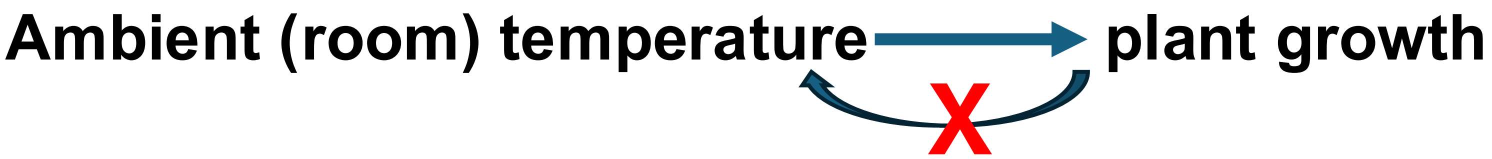

Ambient (room) temperature

plant growth

X

A common point of confusion: Pearson’s correlation is the slope of the best-fit line ***after standardizing* both** variables. The regression slope is the **unstandardized** version that tells you how much Y changes per unit X.

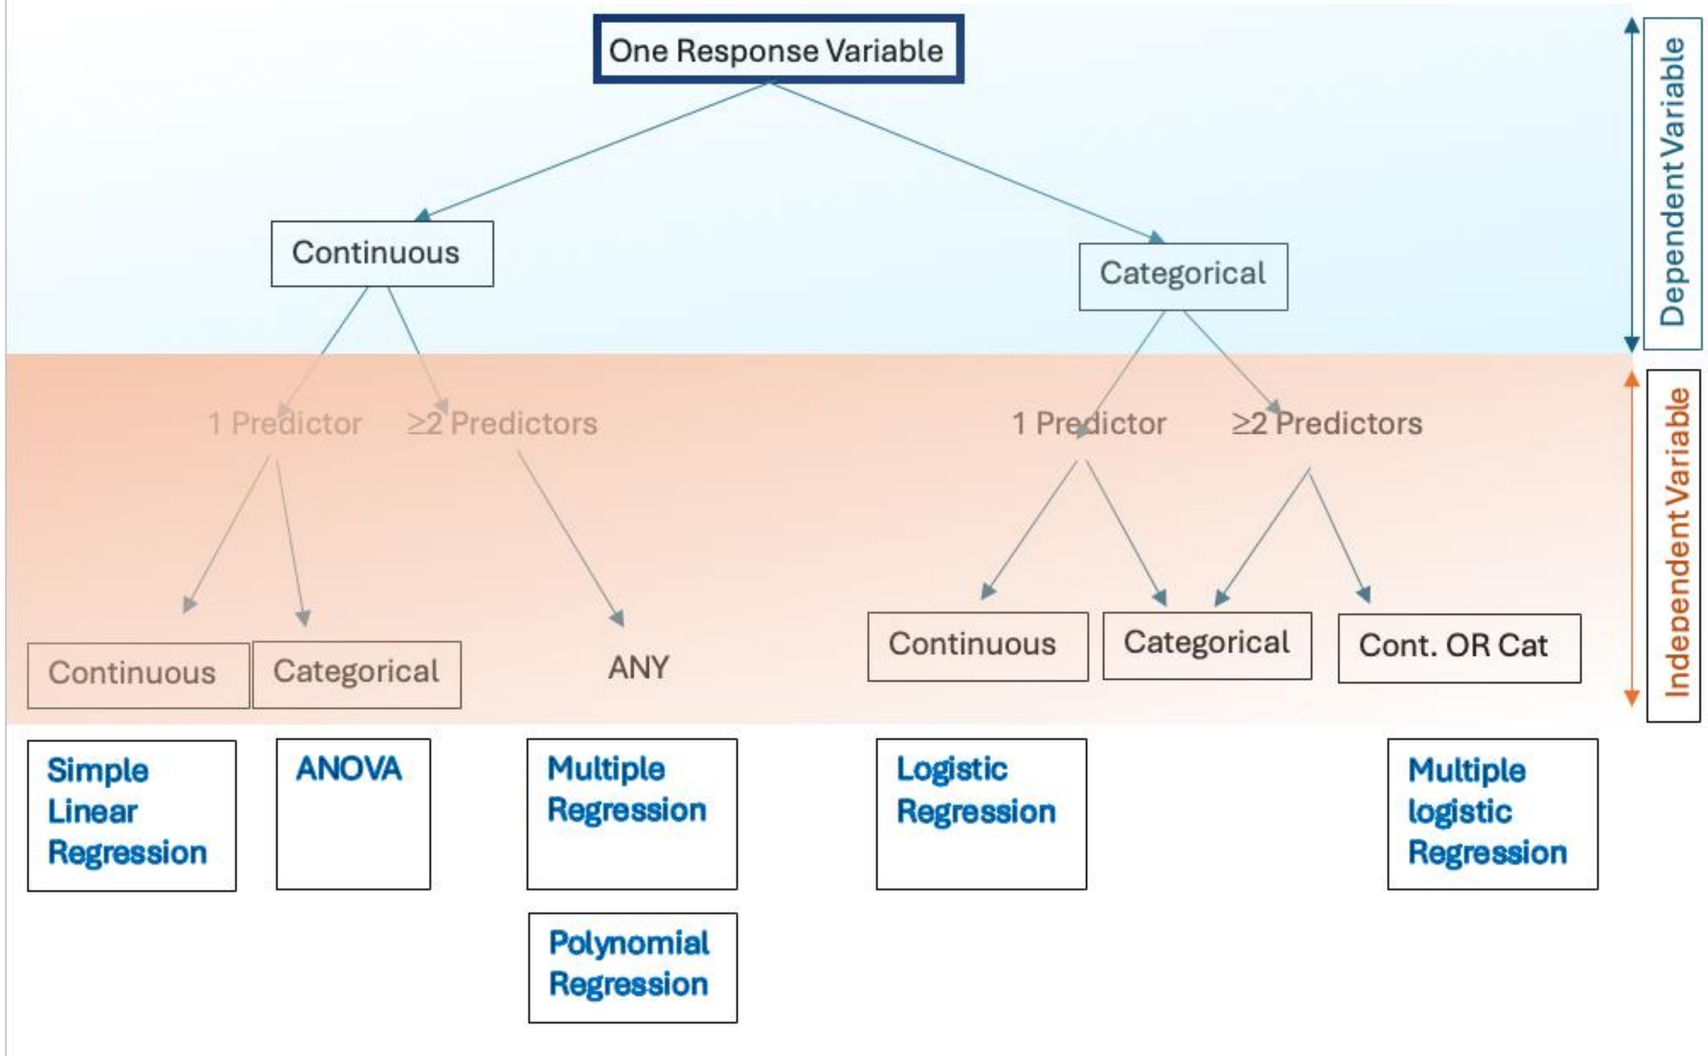

One Response Variable

Continuous

Categorical

Dependent Variable

1 Predictor

≥2 Predictors

1 Predictor

≥2 Predictors

Independent Variable

Continuous

Categorical

ANY

Continuous

Categorical

Cont. OR Cat

Simple Linear Regression

ANOVA

Multiple Regression

Polynomial Regression

Logistic Regression

Multiple logistic Regression

- <u>Other kinds:</u>
  - Lasso
    - Variable selection (weighting a predictor variable by 0)
  - Ridge
    - Allows analysis even in the face of **collinearity**

Regression:

- **Linear** regression assumes that the relationship between X and Y can be described by a line
  - Fits a straight line to a (messy) scatterplot

- **Homoscedasticity**: Y is normally distributed with equal variance for all values of X

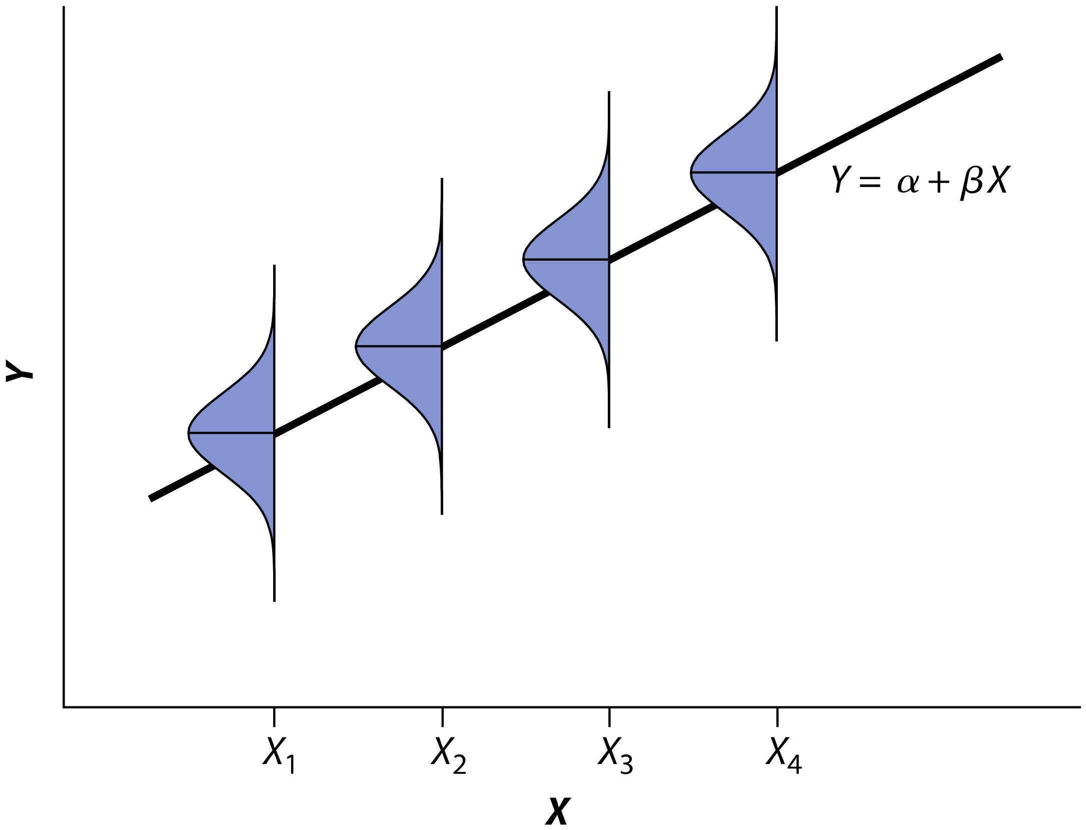

correlation vs regression

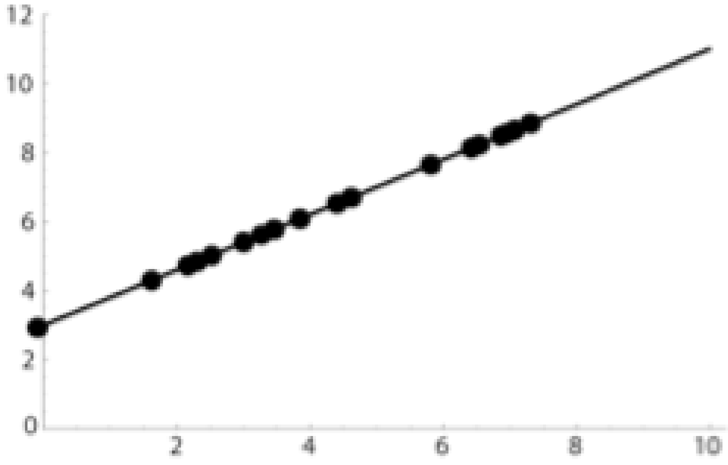

r = 1; Y = 3 + 0.8X

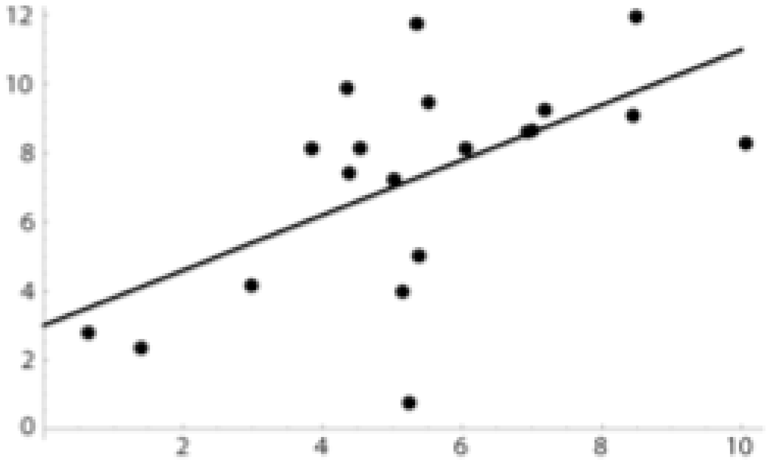

r = 0.6; Y = 3 + 0.8X

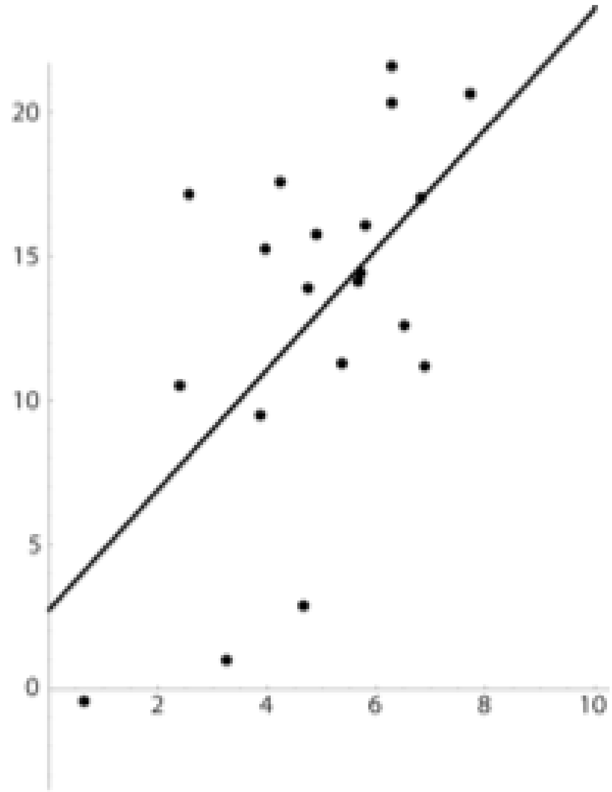

r = 0.6; Y = 3 + 2 X

←

Different correlation;
same slope

→

Same correlation;
different slope

The parameters of linear regression

$$Y = \underbrace{\boxed{\alpha}}_{\text{intercept}} + \underbrace{\beta}_{\text{Slope*}} X + \varepsilon_1$$

* Very similar to Pearson’s correlation (but normalized only wrt to X)

Regression Overview

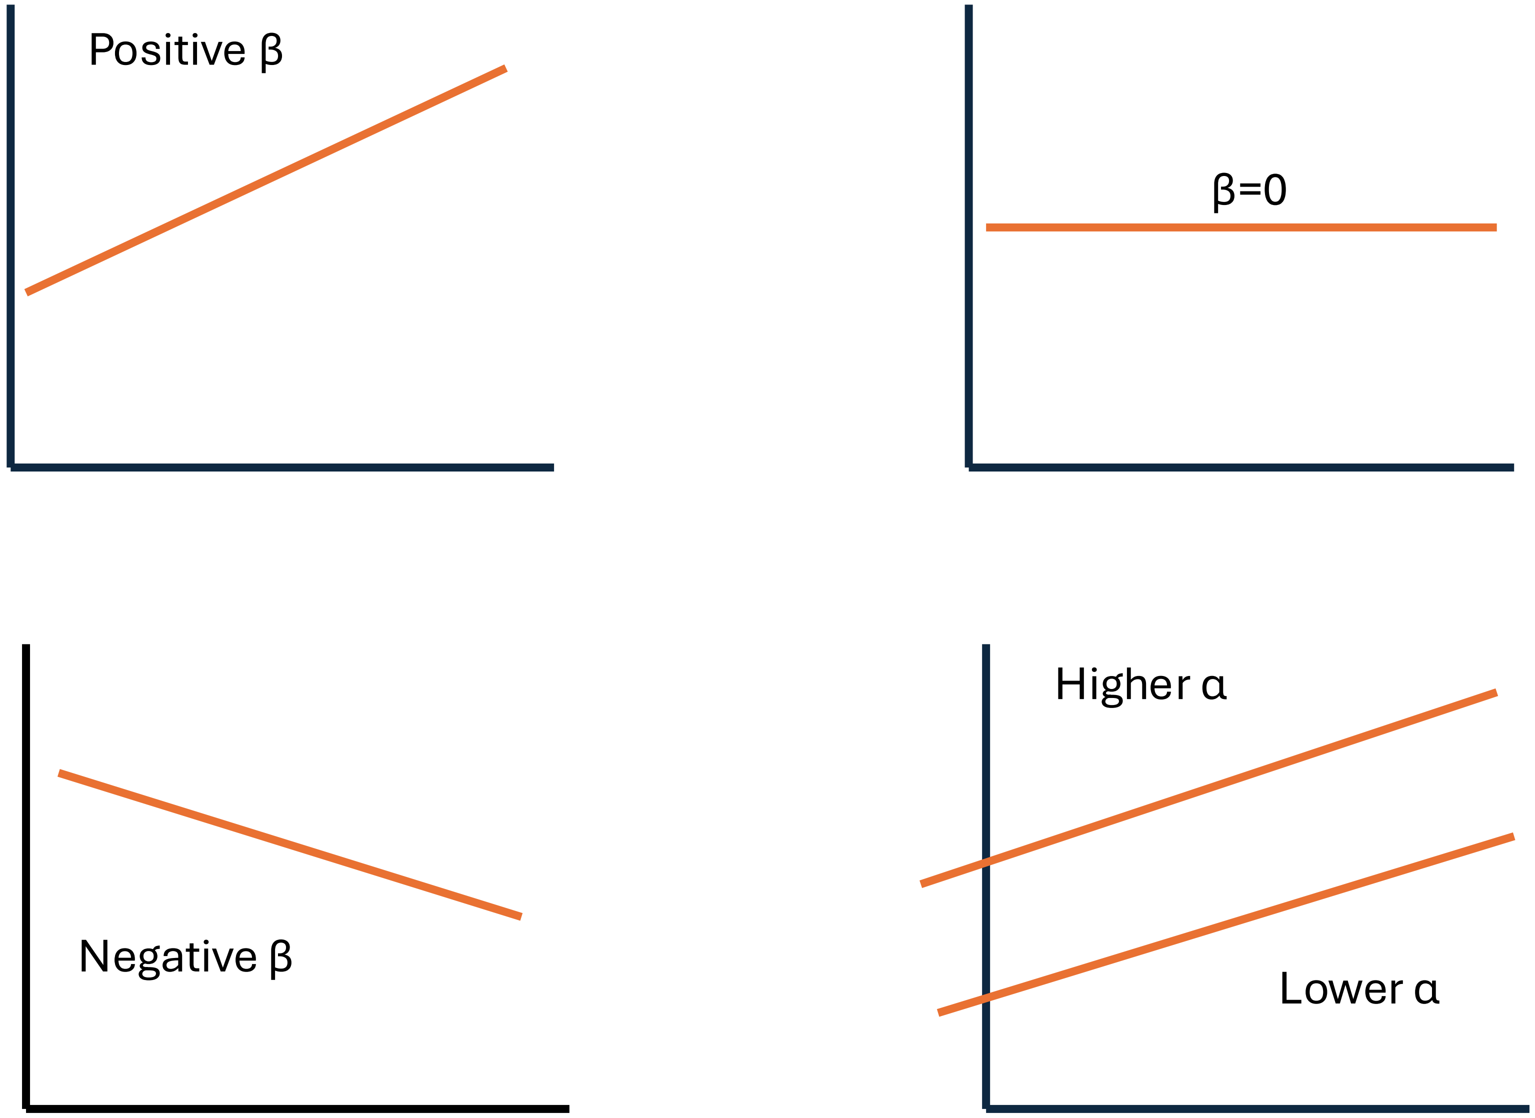

Estimating a regression line

$$Y = a + bX + \varepsilon_1$$

Quick Review:

$$Z = \frac{\bar{Y} - \mu}{\sigma_{\bar{Y}}} = \frac{\boxed{\bar{Y} - \mu}}{\boxed{\sigma/\sqrt{n}}}$$

$$\text{F-value} = \frac{\boxed{MS_{group}}}{\boxed{MS_{error}}}$$

SIGNAL

NOISE

$$t = \frac{\bar{Y} - \mu}{SE_{\bar{Y}}} = \frac{\boxed{\bar{Y} - \mu}}{\boxed{s/\sqrt{n}}}$$

$$r = \frac{\boxed{\square\,(X - \bar{X})(Y - \bar{Y})}}{\boxed{\sqrt{\square\,(X - \bar{X})^2}\sqrt{\square\,(Y - \bar{Y})^2}}} = \frac{Co\,\mathrm{var}\,iance(X,Y)}{s_x s_y}$$

(Ordinary) Least Squares:

- Best fitting line through a scatterplot
  - Line that minimized spread of y values

- Minimize $SS_{residuals}$
  - Measurement of how much the line's predicted y_i deviate from actual data values

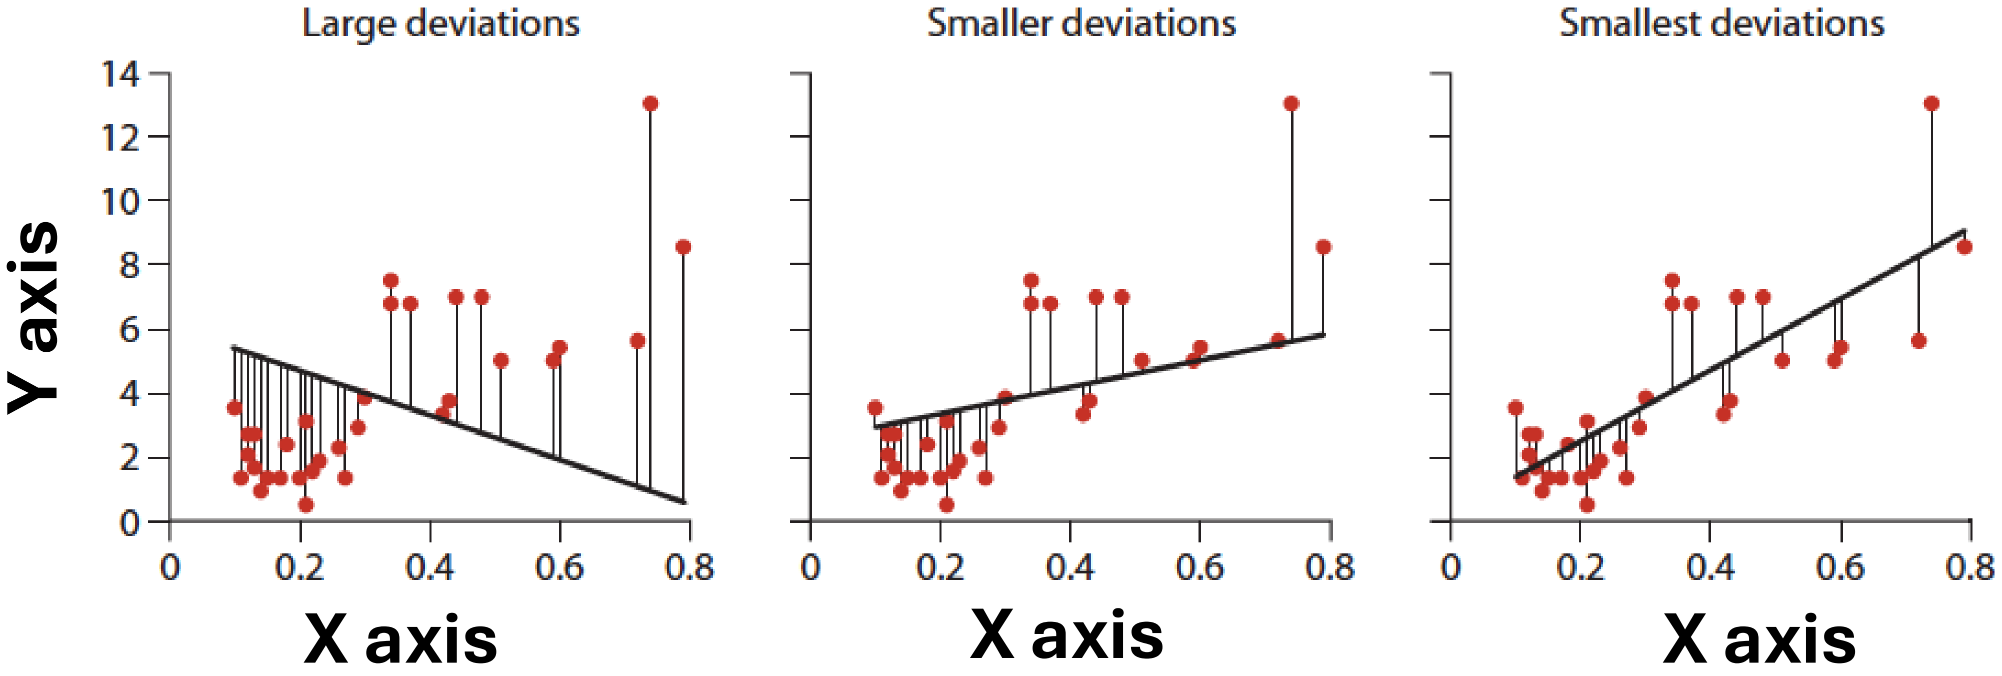

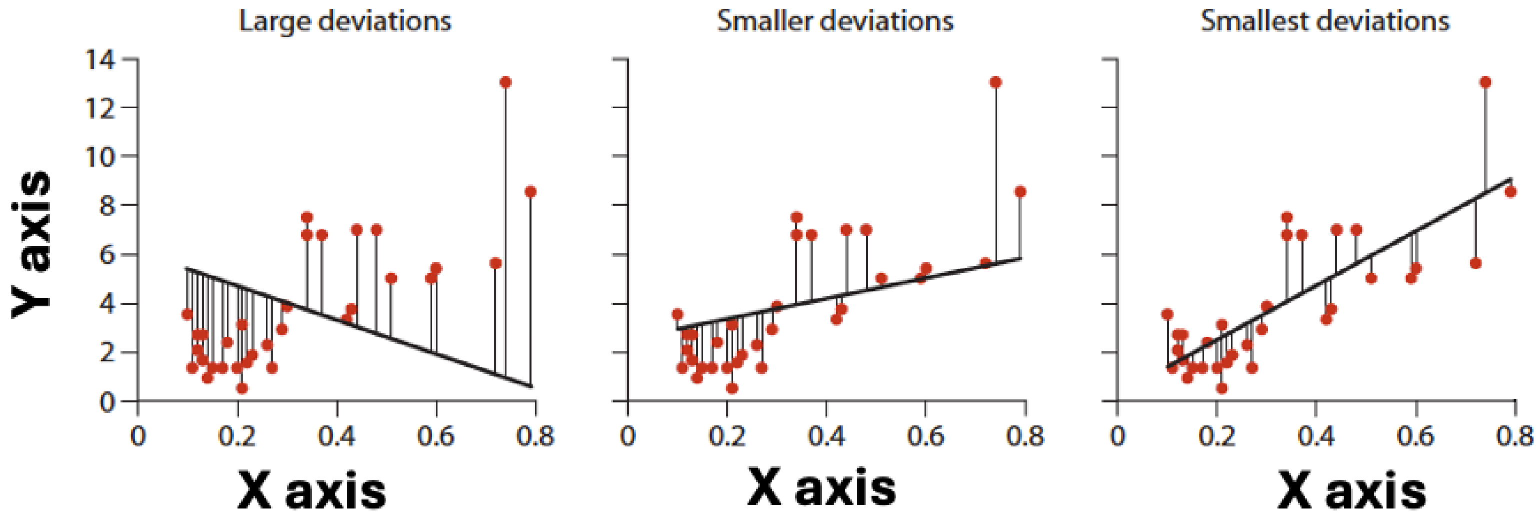

In Linear regression, the line of best fit (OLS) is the GRAND Mean in ANOVA

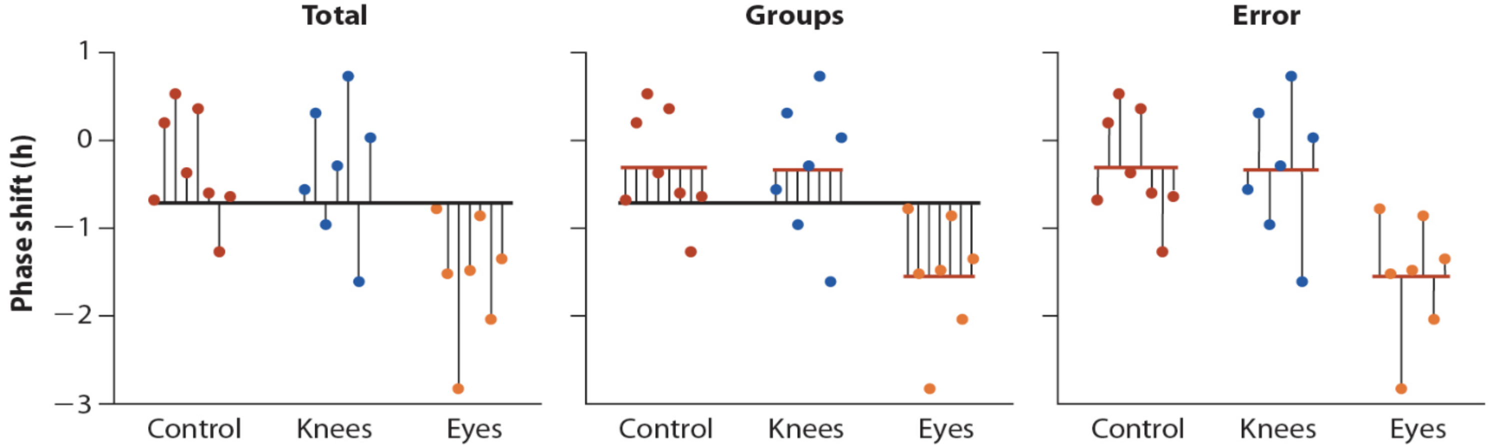

Figure 20.1: Whitock and Schluter, Fig 15.1.2 – Illustrating the partitioning of sum of squares into MS_{group} and MS_{error} components.

Regression Overview

Least Squares:

- What are the elements of this equation?

$$\boxed{SS_{residual} = \sum\limits_{i=1}^{n} \square \, (y_i - \hat{y}_i)^2}$$

$$\boxed{\hat{y}_i = a + bx_i}$$

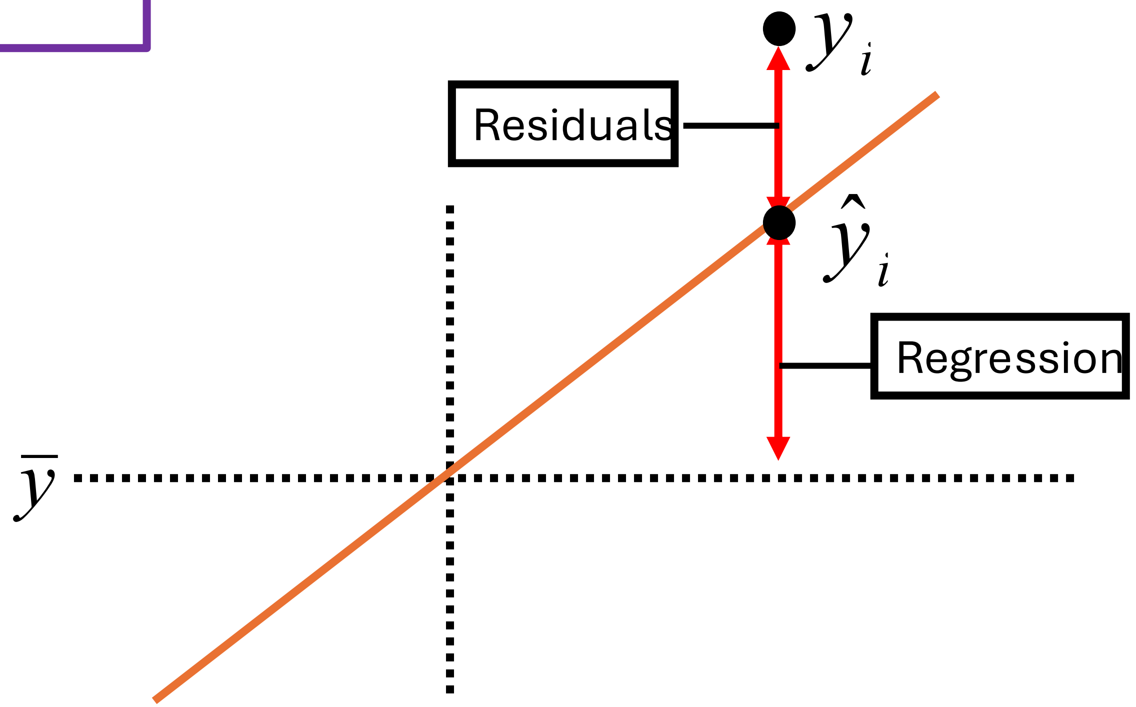

- Residuals:
  - Residuals measure the scatter of points above and below the least squares regression line
  - MS$_{\text{residual}}$ is the variance of the residuals, $residual = Y_i - \hat{Y}_i$

$$\boxed{MS_{residual} = \frac{\square\,(Y_i - \hat{Y}_i)^2}{n-2}}$$

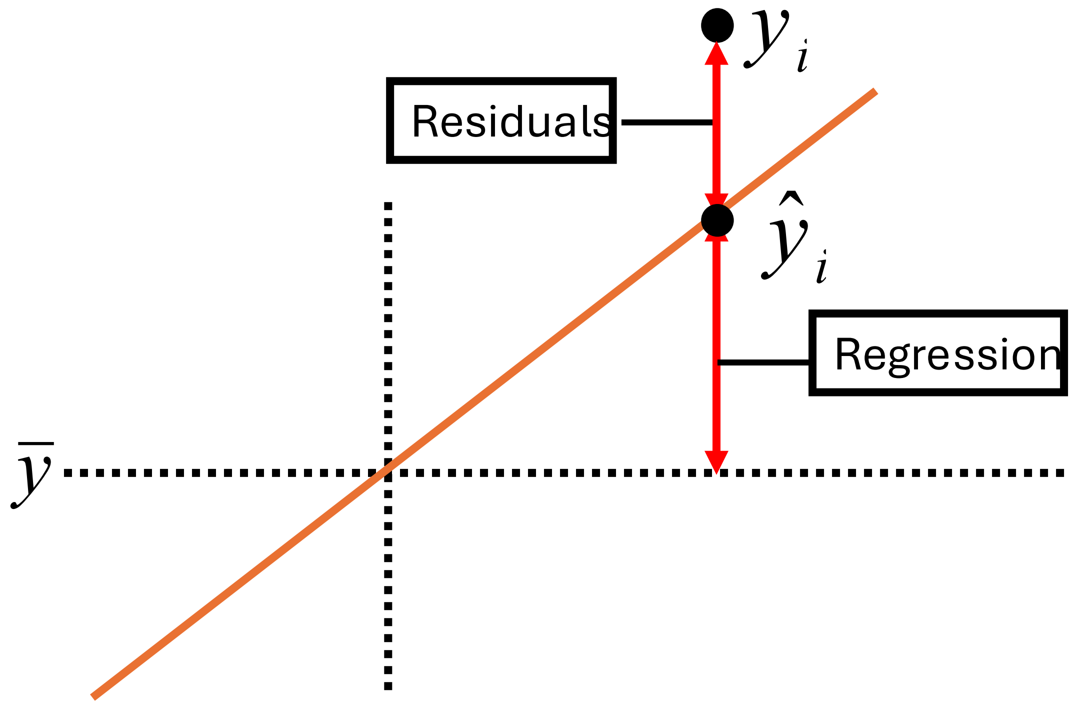

  - MS$_{\text{regression}}$ is the variance of the regression

$$MS_{regression} = \frac{\Sigma(\hat{Y}_i - \bar{Y}\ )^2}{n-2}$$

  - Coefficient of determination (r^2) = SSR/SST

R^2 predicts the amount of variance in Y explained by the regression line

We saw this in ANOVA where R^2 gave ‘precision’ of model (i.e. Ability of the model to explain variation)

- The coefficient of determination
- Sometimes written as r^2
- Square of the correlation coefficient, r

$$R^2 = \frac{SS_{regression}}{SS_{Total}}$$

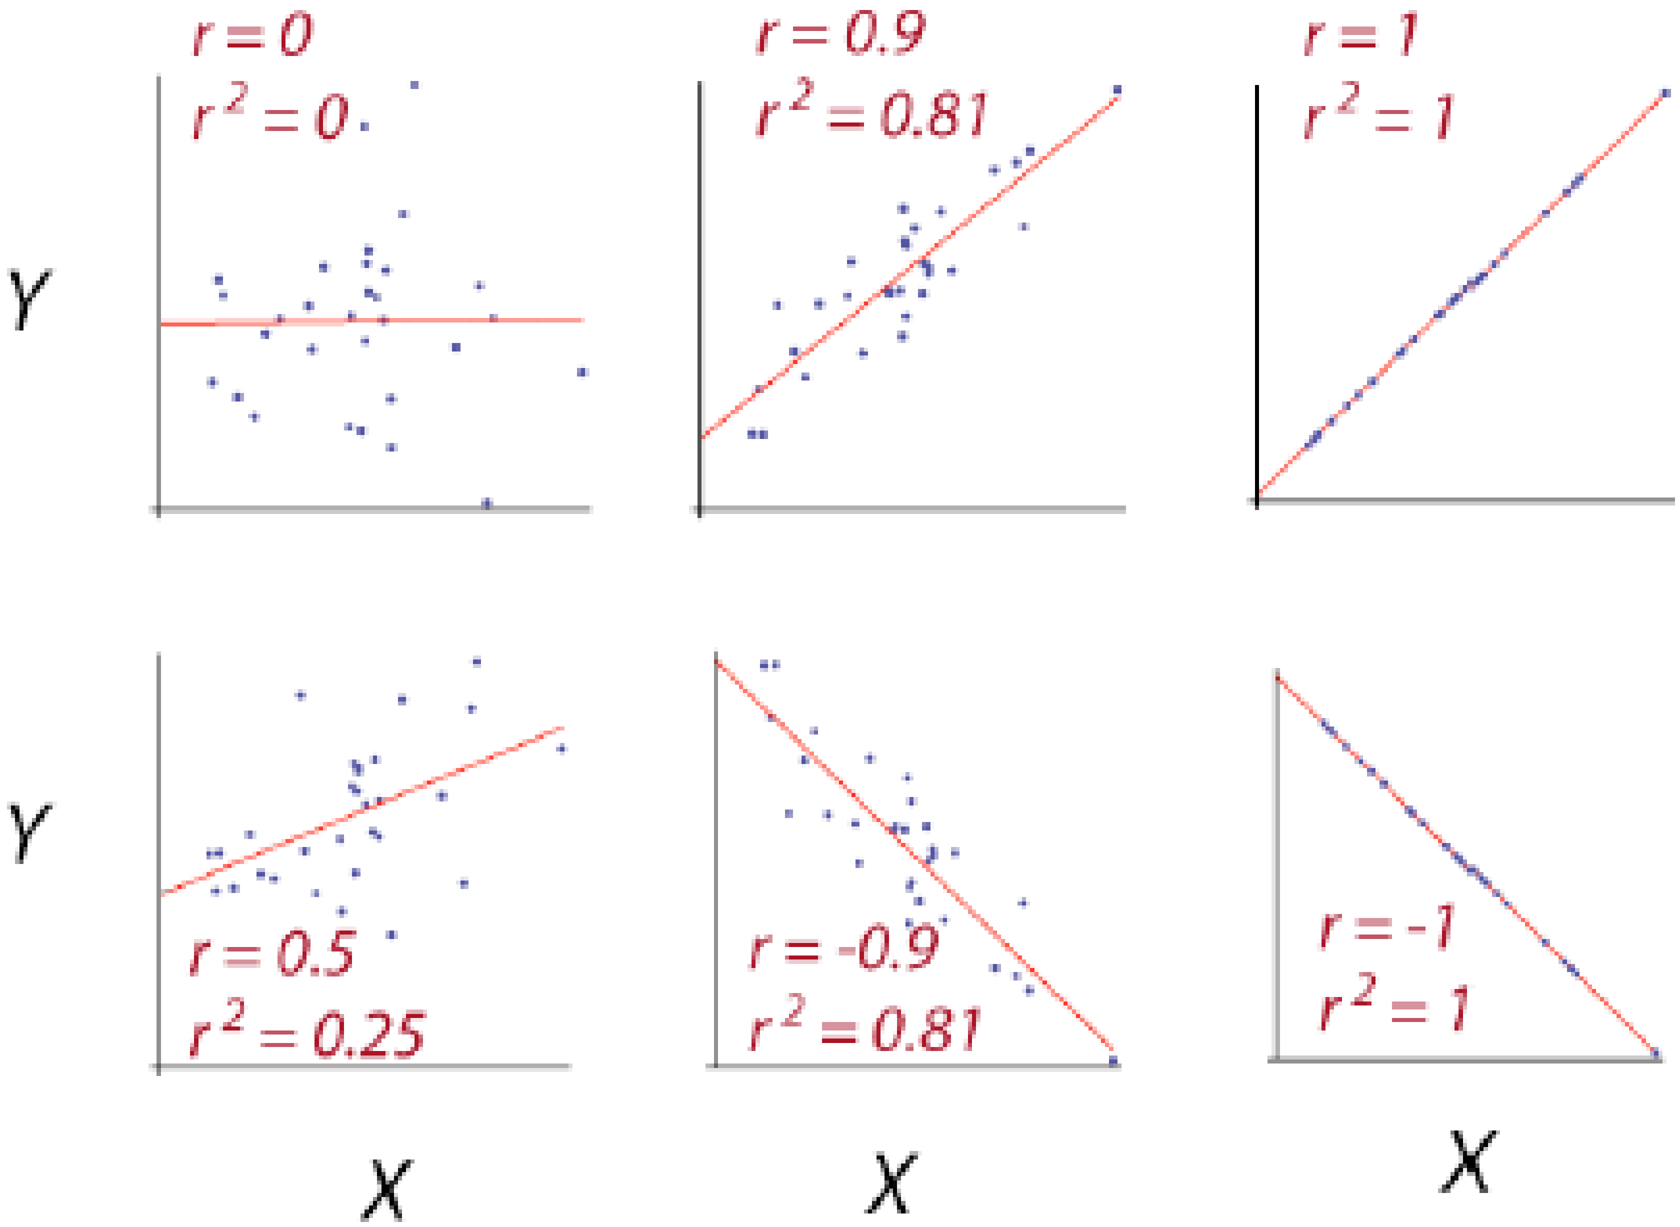
r = 0
r² = 0
r = 0.9
r² = 0.81
r = 1
r² = 1
Y
Y
r = 0.5
r² = 0.25
r = -0.9
r² = 0.81
r = -1
r² = 1
X
X
X

Best estimate of slope:

b = Sum of cross products / Sum of squares of X

$$\boxed{b = \frac{\sum_{i=1}^{n}(X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^{n}(X_i - \bar{X})^2}}$$

Best estimate of slope:

$$b = \frac{\text{Sum of cross products}}{\text{Sum of squares of X}}$$

$$b = \frac{\sum_{i=1}^{n} (X_i - \overline{X})(Y_i - \overline{Y})}{\sum_{i=1}^{n} (X_i - \overline{X})^2}$$

$$r = \frac{\sum (X - \overline{X})(Y - \overline{Y})}{\sqrt{\sum (X - \overline{X})^2}\sqrt{\sum (Y - \overline{Y})^2}} = \frac{Covariance(X,Y)}{s_x s_y}$$

Denominator normalizes based on **both** X and Y variables

Denominator ONLY normalizes based on X (Independent/Explanatory variable)

Finding **a**:

$$\boxed{\bar{Y} = a + b\bar{X}}$$

OR

$$\boxed{a = \bar{Y} - b\bar{X}}$$