

Solving Problems

Effectively collaborating with Ai

Day 1: Knowledge, Uncertainty, and Calibration


Crisis of Authority & Metacognition

Why do smart people make mistakes?

The Calibration Problem: How do we know when to trust our own conclusions?

Course Introduction & Framing

The Calibration Problem: How do we know when to trust our own conclusions?



$$\text{Ai} = \text{Data} \times (\text{Representation} \Rightarrow \text{Loss}) \times \text{Optimization}$$

Workshop Overview: *Why do smart people make predictable mistakes?*

Day 1- What is Knowledge?

Discussion: the crisis of authority and need for metacognition
What is knowledge?
Kahneman systems
Intro to 4 frameworks

Day 2 – What is Reasoning?

Discussion: CRT & Paper
Structure of an Argument
Paul-Elder Framework
Intellectual Autopsy
Polya's 4 steps

Day 3 – Data Literacy & Decisions

Discussion: Dataset
Epistemology of Data
Wild & Pfannkuch
PPDAC Cycle
Synthesis of the four frameworks
Turing Audit

The Calibration Problem: How do we know when to trust our own conclusions?

Thinker's Log

Entry 1

Thinker's Log Entry #1

- **Prompt:** What did I believe confidently this week that turned out to be wrong or uncertain?

The Crisis of Authority

Discussion

Discussion: The Crisis of Authority

Name a time an expert was confidently, publicly wrong and it mattered

Discussion: The Crisis of Authority

Name a time an expert was confidently, publicly wrong and it mattered

1. How do we know what we know, and how do we catch ourselves when we're wrong?
2. If peer review is broken and Ai hallucinates, how can we identify the truth?

If smart people are sometimes wrong, how can we trust our own conclusions?

What is Knowledge, and Why is it “breaking”?

Lecture

Disease

- A disease affects 1 in 1,000 people. A test for it is 99% accurate, meaning that it correctly identifies 99% of people who have it, and correctly clears 99% of people who don't.
- You test positive. What is the probability that you actually have the disease

Disease De-Brief

- **Did the answer feel surprising to you?**
- Three types of Uncertainty
 1. **Aleatoric Uncertainty**
 - Randomness; some things cannot be known
 2. **Epistemic Uncertainty**
 - Missing knowledge; can be reduced using **Ai**
 3. **Model Uncertainty**
 - Representation may be inaccurate; can be increased by using **Ai**

Disease De-Brief

- Did the answer feel surprising to you?
- If new evidence arrives, how should a rational thinker change their belief?
- A positive test increases our confidence but is still not **certain**.

Disease De-Brief

- Did the answer feel surprising to you?
- If new evidence arrives, how should a rational thinker change their belief?

Prior Belief



New Evidence



Updated Belief

Disease De-Brief

- Did the answer feel surprising to you?
- If new evidence arrives, how should a rational thinker change their belief?

Prior Belief



New Evidence



Updated Belief

The **PRIOR**

$$P[A | B] = \frac{P[A \cap B]}{P[B]} = \frac{P[A]P[B | A]}{P[B]}$$

Updated **Belief**

The new **Evidence**



Disease De-Brief

- A disease affects 1 in 1,000 people. A test for it is 99% accurate, meaning that it correctly identifies 99% of people who have it, and correctly clears 99% of people who don't.
- You test positive. What is the probability that you actually have the disease?

Prior: Disease prevalence = 1/1,000

Evidence:
positive test

$$\text{Updated} = P[D|+ \text{ test}] = \frac{0.001 * 0.99}{(0.001 * 0.99 + 0.999 * 0.01)}$$

1. $P[+ \text{ive test} | \text{Disease}] = 0.99$
2. $P[+ \text{ive test} | \text{No Disease}] = 0.01$

negative test

1. $P[- \text{ive test} | \text{Disease}] = 0.01$
2. $P[- \text{ive test} | \text{No Disease}] = 0.99$

Disease De-Brief

- A disease affects 1 in 1,000 people. A test for it is 99% accurate, meaning that it correctly identifies 99% of people who have it, and correctly clears 99% of people who don't.
- You test positive. What is the probability that you actually have the disease?


Prior: Disease prevalence = 1/1,000

Evidence:

positive test

1. $P[\text{+ive test} | \text{Disease}] = 0.99$
2. $P[\text{+ive test} | \text{No Disease}] = 0.01$

$$P[D | \text{+ test}] = \frac{0.001 * 0.99}{(0.001 * 0.99 + 0.999 * 0.01)}$$



Disease De-Brief


- A disease affects 1 in 1,000 people. A test for it is 99% accurate, meaning that it correctly identifies 99% of people who have it, and correctly clears 99% of people who don't.
- You test positive. What is the probability that you actually have the disease?
- Calibration = appropriate Bayesian updating

When encountering new evidence, you can adopt the following approaches:

1. Ignore new evidence
 2. Over-updating
 3. Appropriate updating
- Miscalibration = failure to update appropriately
 - Important question before looking at evidence:

What evidence would change your belief?

What is knowledge and why is it breaking?



Plato's Cave


[wikipedia](https://en.wikipedia.org/wiki/Plato's_Cave)

What is knowledge and why is it breaking?

- Knowledge is justified true belief — Plato (428 - 348 BC)
 - Justification
 - Method choice, model selection
 - Truth
 - Honesty, transparency
 - Belief
 - Prior informative or experience, cognitive biases
- If knowledge is provisional, how can we rationally update beliefs in the light of new data?

What is knowledge and why is it breaking?

Knowledge NEVER speaks for itself




- Difference between knowing THAT and knowing HOW
- We're going to examine a few frameworks for knowing

What is knowledge and why is it breaking?

- The Gettier Problem (1963. Only 3 pages!)

You can have justified true belief that is true *by accident*.

Is it still considered knowledge?



What is knowledge and why is it breaking?

- **Tacit Knowing – Michael Polyani (1966)**

- “We can know more than we can tell”
- Knowledge is unarticulated, and intuitive (i.e. riding a bike)
- Knowledge is subjective
- Polanyi’s paradox: The difficulty Ai systems have in replicating knowledge that humans cannot explicitly articulate.

What is knowledge and why is it breaking?

- **Social Epistemology**

- Knowledge is not produced individually; it is produced in communities with norms
- Some voices are amplified
- **Epistemic injustice - Miranda Fricker(2007)**

What is knowledge and why is it breaking?

There is a crisis of epistemic authority

- A collapse of trust in experts, institutions, and credentialed knowledge.
- Frankfurt (2005): *On Bullshit*
 - Gives you the necessary vocabulary
- Nichols (2017): The Death of Expertise
 - The internet did not democratize knowledge, it democratized ***the feeling of being informed***
 - This is more dangerous than being and knowing that you are uninformed

What is knowledge and why is it breaking?

There is a crisis of epistemic authority

AI gives everyone a fluent, plausible, and confident *expert voice* **at scale**

- *Crockett & Messeri (2024, Nature): The Illusion of Understanding*. We outsource our thinking, retain the confidence, and lose the ability to detect when we are wrong. This is Gettier at scale!
- LLMs produce Bullshit, which is **distinct from lying**.
- *convincingness is uncorrelated with accuracy*

What is knowledge and why is it breaking?

There is a crisis of epistemic authority

AI gives everyone a fluent, plausible, and confident *expert voice* at scale

We need to build structural tools to combat this!

- *We cannot trust authority*
- *We cannot even trust ourselves*



System 1/ System 2

Lecture

System1/System 2 – Daniel Kahneman

1. A bat and a ball cost \$1.10 in total. The bat costs \$1 more than the ball. How much does the ball cost?

System1/System 2 – Daniel Kahneman

1. A bat and a ball cost \$1.10 in total. The bat costs \$1 more than the ball. How much does the ball cost?
2. If it takes 5 machine 5 minutes to 5 widgets, how long for 100 machines to make 100 widgets?

System1/System 2 – Daniel Kahneman

1. A bat and a ball cost \$1.10 in total. The bat costs \$1 more than the ball. How much does the ball cost?
2. If it takes 5 machine 5 minutes to 5 widgets, how long for 100 machines to make 100 widgets?
3. In a lake, there is a patch of lily pads. Every day, the patch doubles in size. If it takes **48** days for the patch to cover the entire lake, how long would it take for the patch to cover half the lake?


System1/System 2 – Daniel Kahneman

Correct answers to the questions:

System1/System 2 – Daniel Kahneman


System 1 – Fast Thinking

What answer **feels** right?



System1/System 2 – Daniel Kahneman

System 2- Slow thinking



System1/System 2 – Daniel Kahneman

Characteristics

System 1 - FAST	System 2 - SLOW
Automatic	Deliberate
Effortless	Effortful
Pattern-Based	Analytical, Logical, Sequential

Used For

System 1 - FAST	System 2 - SLOW
IMMEDIATE reactions	Complex decisions
Reading facial expressions	Planning
Expert intuition in familiar domains	Mathematics, statistical reasoning, evidence evaluation

System1/System 2 – Daniel Kahneman

Characteristics

System 1 - FAST	System 2 - SLOW
Automatic	Deliberate
Effortless	Effortful
Pattern-Based	Analytical, Logical, Sequential

- Relies on shortcuts
- Vulnerable to bias
- Overconfident
- Substitutes easy questions for hard ones

- Mentally expensive
- Avoided unless necessary

Used For


System 1 - FAST	System 2 - SLOW
IMMEDIATE reactions	Complex decisions
Reading facial expressions	Planning
Expert intuition in familiar domains	Mathematics, statistical reasoning, evidence evaluation



System1/System 2 – Daniel Kahneman

Cognitive Biases

1. Confirmation bias
2. Availability Heuristic
3. Anchoring
4. Overconfidence
5. Motivated Reasoning
6. Others*



System1/System 2 – Daniel Kahneman

What is more dangerous in the US: Sharks or Vending Machines?





System1/System 2 – Daniel Kahneman

Strategies to Engage system 2?

Good problem solving is not about eliminating intuition.

It is about recognizing when intuition is insufficient.

Frameworks

Four frameworks designed to catch failures



The Four Frameworks

The Calibration Problem

How do we know when to trust our own conclusions?

Framework	The Calibration Failure	The Mechanism	Day
Kahneman System 1/System 2	We don't recognize fast thinking	Name the system; interrupt the response	1
Paul-Elder Critical Thinking	We argue from hidden assumptions that are never examined	Expose the structure; demand evidence for each element	2
Polya Problem Solving	We treat complex problem as unsolvable wholes	Decompose; transform ambiguity into tractable steps	2
Wild & Pfannkuch Statistical Thinking	We trust data that has errors we cannot see	Interrogate the measurement before trusting the conclusion	3

The Four Frameworks

Kahneman

System 1 - *FAST*

- Uses heuristics
- Efficient
- Biases and errors

System 2 – *SLOW*

- Deliberate, analytical
- Requires attention
- Working memory
- More accurate
- Mentally demanding

”Good problem solving is not about eliminating intuition – it’s about knowing when intuition is insufficient – Daniel Kahneman, *Thinking, Fast and Slow* (2011)

The Four Frameworks

Polya

1. Understand Problem

- Clarify
- Identify knowns, unknowns, constraints
- Visualizations

2. Devise a Plan

- Choose strategy
- Break problem into smaller parts
- Similarity to other easier (already-solved) problems

3. Execute Plan

- Execute strategy
- Check each step

4. Iteration & Reflection

- Verify solution, “does it make sense?”
- Generalize lessons
- Consider alternatives

“A good solver is not someone who immediately knows the answer, but someone who systematically explores the problem” – George Polya, *How to Solve it* (1945)



The Four Frameworks

Paul-Elder

1. Purpose

- What is the goal?

2. Question at Issue

- What problem or question are we addressing?

3. Information

- What evidence or data do we have?

4. Interpretation & Inference

- What conclusions are we drawing?

5. Concepts

- What theories, definitions, or principles are involved?

6. Assumptions

- What are we taking for granted?

7. Implications & Consequences

- What follows from this reasoning?

8. Point of View

- What perspectives are influencing the analysis?

“Critical thinking improves when we actively examine the quality of our reasoning.”

Richard Paul & Linda Elder, *The Miniature Guide to Critical Thinking*



The Four Frameworks

Wild & Pfannkuch

1. Problem

- Define the question
- Clarify context & goals

2. Plan

- Decide on what data is needed
- Design methods

3. Data

- Collect, organize and clean data
- Consider data quality and limitations

4. Analysis

- Explore patterns, relationships, and variation
- Use graphs, models, and summaries

5. Conclusion

- Interpret findings in context
- Evaluate uncertainty and limitations
- Communicate results appropriately

“Data alone do not produce insight - insight comes from reasoning about variation, context, and uncertainty.”

- Wild & Pfannkuch, “*Statistical Thinking in Empirical Enquiry*” (1999)



Hiring Game

Activity

Hiring Bias Game #2

1. Design rubric
2. Assess **6** applicants
3. Discuss your criteria with the group

Hiring Game De-Brief

- How well does your rubric fully explain your decisions?
- Is there still a gap between your stated rubric and the actual operating assumptions?
- What is your rubric?

The Three Questions

- Confidence should track evidence

Confidence	Accuracy
95%	?
80%	?
60%	?

- What would a well-calibrated thinker look like?

The Three Questions

1. What do I know?
2. What don't I know?
3. How much does that ignorance matter?

The Four Frameworks

Kahneman

System 1 - *FAST*

- Uses heuristics

System 2 – *SLOW*

- Deliberate, and analytical
- Cognitively demanding

Polya

1. Understand Problem
2. Devise a Plan
3. Execute Plan
4. Iteration & Reflection

Paul-Elder

1. Purpose
2. Question at Issue
3. Information
4. Interpretation & Inference
5. Concepts
6. Assumptions
7. Implications & Consequences
8. Point of View

Wild & Pfannkuch

1. Problem
2. Plan
3. Data
4. Analysis
5. Conclusion

The Four Frameworks

Kahneman

System 1 - *FAST*

- Uses heuristics

System 2 – *SLOW*

- Deliberate, and analytical
- Cognitively demanding

Polya

1. Understand Problem
2. Devise a Plan
3. Execute Plan
4. Iteration & Reflection

Paul-Elder

1. Purpose
2. Question at Issue
3. Information
4. Interpretation & Inference
5. Concepts
6. Assumptions
7. Implications & Consequences
8. Point of View

Wild & Pfannkuch

1. Problem
2. Plan
3. Data
4. Analysis
5. Conclusion

Calibration Failures these model's address:

We are unaware that fast thinking is operating at all!

We treat complex problems as unsolvable wholes!

We argue from hidden assumptions that we never examine!

We trust data that encodes errors that we cannot see!

Homework