Crowd sourcing popular 'AI' (broadly defined) books via JAX

	Book	Summary
1	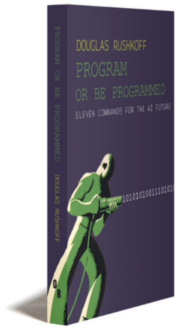	A slim book with a blunt message: either we shape technology, or it shapes us. The "commands" are about literacy - how to keep from being programmed by the very tools we've built. I treat technology as if it is neutral, when it's anything but. There is an unsettling sense that every click, scroll, or "like" is part of someone else's program, unless I resist.
2	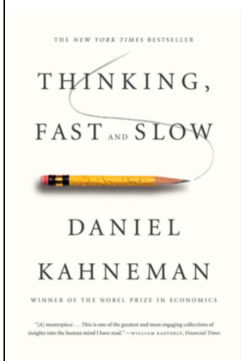	A classic that is about, well, thinking. Premise is that we have two systems of thinking. After years of teaching statistics, I already had a prior belief that our brain rebels against learning certain ideas (like statistics). I initially read this book because I realized that I did not understand what 'learning' is. I am still not sure that I do, but I am a more informed kind of ignorant.
3	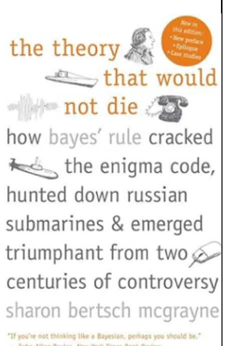	I will fight anyone who doesn't recognize the supremacy of Bayes' theorem in Bioinformatics etc. This is basically the biography of Bayes' theorem - which sounds dull, but it's weirdly engaging. The story winds from obscure clergymen to wartime codebreakers to modern data science, and along the way makes clear that probability isn't just math but a worldview. It's a reminder that even equations can be political – which is a theme for many of these broadly data science books.
4	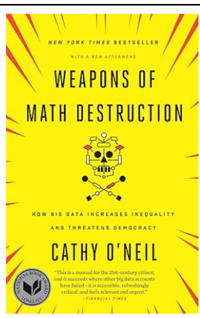	O'Neil takes a wrecking ball to the idea that algorithms are impartial. Her "weapons of math destruction" are the models that scale inequality while hiding behind a veneer of objectivity. I found myself nodding a lot since I tend to emphasize assumptions and limitations of models. Note: She has a fantastic blog called "mathbabe.org"

5	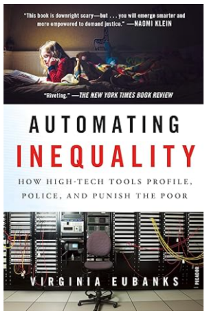	Eubanks shows how supposedly neutral systems, welfare eligibility software, risk assessments, automated decision-making, recreate poverty traps with a gloss of “objective authority”. As geneticist, there will be some echoes of historic problematic eugenics-adjacent ideas that were supposedly grounded in ‘objective’ facts in genetics/biology. The book hits harder than most academic critiques because the case studies are so concrete, so numerous, and so brutal. I came away thinking less about technology itself and more about power: who designs, who benefits, who suffers consequently. This book confirmed suspicions I’ve carried for years: sometimes an unnecessary tech “solution” serves as an excuse to avoid human responsibility. There is a warning about just following orders in this book. This book is *dense*, since it is full of case studies and, because of that, despite illustrating persistent, evergreen issues, it is not necessarily au courant (that is: there will be newer examples that will illustrate the same concepts).
6	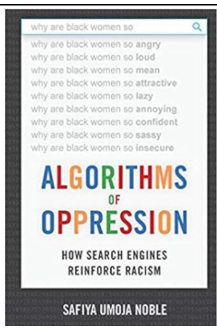	Search engines – an accessible example of large data for almost everyone in the world – doesn’t just “reflect” the world. Noble shows how they encode and amplify racist and sexist biases. I teach about bias in data, but seeing it spelled out in the everyday act of Googling was sobering. It left me with the uncomfortable realization that even my most mundane digital habits are implicated in structural injustice.
7	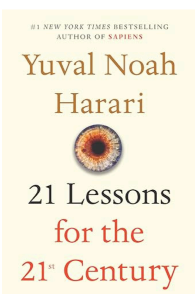	Harari (Aka the new “Malcolm Gladwell”) synthesizes lots of pieces into grand pronouncements, some of which age better than others. This one felt less like a roadmap and more like a meditation on how overwhelming it is to live now, including AI, nationalism, climate, religion, everything all at once. I didn’t agree with every “lesson,” but that almost felt like the point. It’s the kind of book that leaves you anxious thinking about what can go wrong, and how tools can be manipulated in scalable ways… which a lot of this list focuses on. This is for a popular audience, not a technical audience.
8	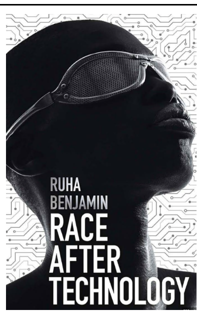	Introduces the term: “New Jim Code,” and shows how racism is coded into the very architecture of emerging technologies. It’s not just biased outcomes, though, it’s the design choices, defaults, and examples that reproduce inequality. I came in to the book expecting another “algorithms are biased” argument (which, fair enough, they are and that point needs to emphasized all the time!), but her framing shifted my understanding from outcomes to “systems thinking”. It’s the rare book that both sharpens your analysis and makes you uncomfortable in the best way.
9		Mitchell’s book is the one I wish more AI conversations started with: a measured, clear-eyed, skeptical without being dismissive. She lays out what AI can do, what it can’t, and

	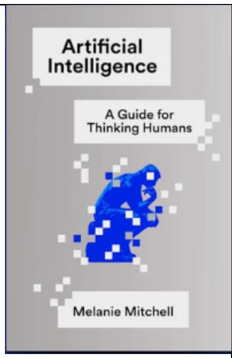	why the hype cycle clouds both categories. Reading it felt like a palate cleanser after wading through headlines and self-serving white papers with grandiose claims. It reminded me that being a “thinking human” about AI mostly means resisting the urge to be impressed too quickly.
10	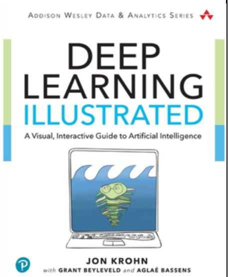	This book is part textbook, part comic book, part … pep talk. It makes neural networks feel approachable, even dare I say, charming, without hand-waving the math too much. I read it thinking it might be too “entry level,” but it helped me explain concepts better. Sometimes the best way to learn something advanced is to see it drawn with bright colors and metaphors instead of Greek letters. Cue obligatory Feynman quote about explaining complicated ideas to a 5 year as a measure of if you understand an idea yourself.
11	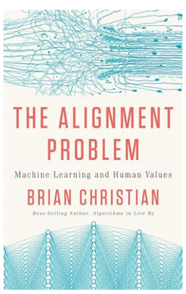	This book weaves technical debates about AI safety with human stories focused on researchers, ethicists, even philosophers wrestling with what it means to align machines with values. What struck me most was how slippery the word “alignment” becomes once you stop pretending human values are tidy categories. Like how many of these books prompt you to stop thinking that algorithms are ‘neutral’ and ‘objective’? I came away less confident that AI can be “aligned” in any absolute sense, but more convinced that the conversation matters anyway (can we align our own behaviour? Can human behaviour and human systems ever actually be aligned?). It’s a book that leaves you with “productive unease”. Yes, this is a pattern for many of these books.
12	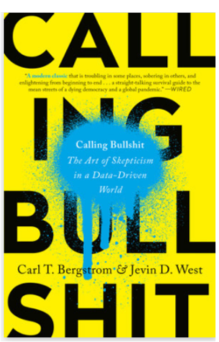	A handbook for spotting nonsense in the age of data, charts, and dashboards. It’s funny, sharp, and occasionally brutal. I stole (with credit, of course) many examples for my various statistics courses. The two authors have a short free course available online that is a companion to this book (but possibly a bit outdated). I felt equal parts entertained and chastised; it’s easy to laugh at others’ mistakes until you realize you’ve made the same ones, despite knowing better! Bullsh!t is a collective problem. (Bullsh!t is not the same thing as lying; lying means that you know what truth is and are purposely being orthogonal to it; Bullsh!t is intended to persuade without any lode star of truth. https://callingbullshit.org/ https://en.wikipedia.org/wiki/On_Bullshit <- description of the original philosophy of Harry Frankfurt.
13		

	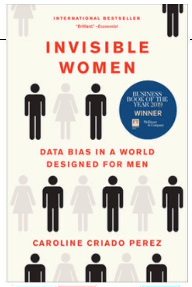	This one is a catalog of absence in data (which is famously difficult to do: "absence of evidence is not evidence of absence): all the ways women are left out of data, design, and decision-making. From car safety to medical trials, the omissions pile up until it's impossible to see them as accidents. I thought I already knew the "gender data gap," but the sheer accumulation outlined in this book was overwhelming. It made me rethink what counts as "neutral" in data science, and how often neutrality is just code for ignoring data (even >50% of the world!).
14	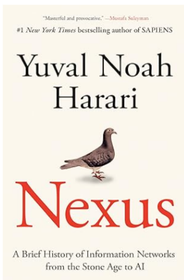	A curated bundle of essays that orbits the question of how humans and technology are merging. It's uneven, contains a lot of speculation, but the focus is his fascination with what happens when biology and data fuse which, of course, is highly relevant to my own interests! It's typical for Harari, and I was impressed by the sweep, but skeptical about the pretense of certainty. It's provocative, though, which I almost always enjoy.
15		This maps the hidden costs of AI that are invisible to me and, probably, other people. This includes the mining, labor, environmental damage, and geopolitical struggles. Some of these are sunk costs (pre-training costs), and some are prescient of what Data Centers might include. What I found most compelling was how unglamorous her picture of AI is: less sleek robots, more warehouses, lithium, and underpaid and exploited (especially the mental toil it takes to curate unimaginably horrific images via data annotation) workers. It made me realize that "intelligence" is often built on the backs of invisible infrastructures. A sobering antidote to AI hype.
16	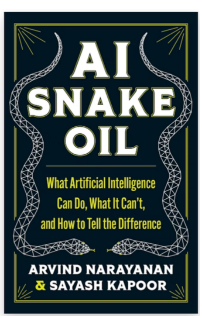	Not everything branded "AI" is "intelligent", useful, or safe (and maybe anything with that label should immediately prompt skepticism). The authors do a great job of separating what works (and is boringly reliable) from what's speculative at best, fraudulent at worst. It balances skepticism without being anti-technology. It's also a reminder that calling out hype is a form of public service, just like in "calling bullsh!t". This book was a bit thin, though, and, while accessible, I am not convinced it offered much new information from the other books on this list (including recycling some prominent examples).
17		This is a bit of an 'inside-baseball' burn book that uses the careers of the leading figures involved in OpenAI (mostly) to illuminate other issues within the AI landscape. It spends a lot of time on the temporary ouster of Sam Altman from OpenAI and uses its dissection and motivations to uncover more hidden issues. There is a lot about the politics of LLMs and "AI" and who *should* have control over such a powerful technology (should any one

	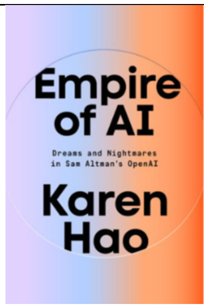	figure?). It was interesting and focused on the Machiavellian nature of this technology. It WILL reshape the world, and so the politics and motivations of the people shaping its zeitgeist matters.
18	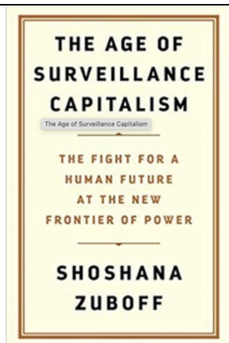	This book argues that major technology companies turned human behavior into a raw material (remember the adage: if you don't pay for a product, you ARE the product…Cough, cough: Facebook): collecting data about what we do, using it to predict what we will do next, and increasingly designing systems that can shape those choices. The central concern is not simply privacy, but power: who gets to observe, predict, and influence behavior, and how little transparency or accountability surrounds that process. It is especially relevant to AI because today's systems depend heavily on data, making questions of consent, ownership, and behavioral influence even more urgent.
19	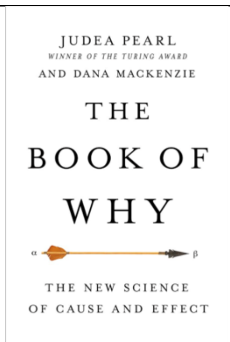	If Machine Learning is just approximately a series of stacked logistic regressions that spits out millions of correlations, what about… causation? This book introduced me to DAGs (Directed Acyclic Graphs) and causal thinking. It has fundamentally shifted my perspective and that is the greatest compliment that I can give a book. I read it four years ago and still quote from it.
20	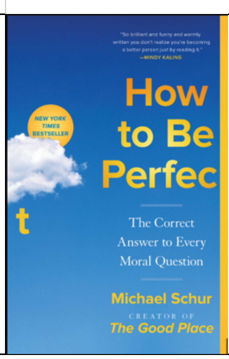	If you are interested in AI, **you are eventually going to need to get caught up on the big ideas of western philosophy.** This book will help you do so, *and* it is funny. Trust me: you don't want to have to read Hume (the man who already thought of all this ~300 years ago), so read this instead. Then *maybe* Hume.

21	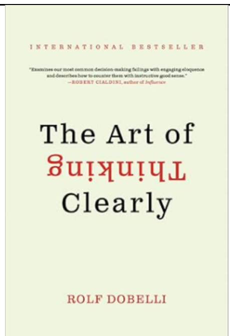	This isn’t a book to read in one sitting, since it is a catalogue of the many biases that we human being have. It is a fun reference, especially to remind yourself of the many judgement errors that you have…even if we like to pretend that we don’t have them.
22	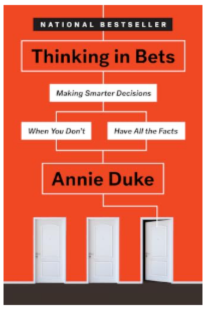	Breaks down decision-making as a process of reasoning under uncertainty. Annie Duke connects game theory, Von Neumann famously treated poker as an important model for strategic decision-making, with Bayesian updating (“hello, Bayes!”), calibration, and the psychology of confidence. One of the book’s most useful ideas is that asking people how much they would actually *bet* on a belief often reveals how confident they really are, rather than how confident they claim to be. Together, these tools offer a practical framework for making better decisions: separate outcomes from decision quality, continually update your beliefs as new information arrives, and get more precise about what you know and how certain you are. The book also covers many of the greatest hits of cognitive bias: how we fool ourselves, how readily we construct stories from incomplete evidence, and why our perception of reality is not necessarily reality itself. Although it was published in 2018 and does not focus on AI, its ideas feel increasingly relevant to current conversations about AI. Both humans and AI systems operate with incomplete information, uncertain predictions, and imperfect representations of the world. Learning to calibrate confidence, update beliefs, and distinguish evidence from a compelling narrative may therefore be just as important for improving how we *interact with AI* as it is for improving our own decision-making.
23	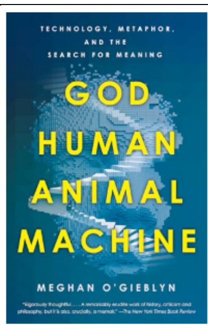	This book explores how AI and digital technology blur the boundaries between humans, animals, machines, and even the divine. The content draws on philosophy, theology, and technology to ask what intelligence, consciousness, and meaning actually are when machines can increasingly imitate human thought. Ultimately, the book is as much about what AI reveals about our assumptions about ourselves as it is about AI itself. It is upsetting, and disquieting. I would put it in the top 3 must-reads from this list.

It’s a bit dangerous to include fictional work in this list, but fools rush in….

1		Besides annoyingly ubiquitous and ridiculous tropes, this book is a prophecy about the internet and virtual reality. It was published in 1992. Ever wonder why (the) Facebook was renamed the Metaverse (Yeah, it came from here… Mark Zuckerberg’s favourite book). Google Earth was, apparently, inspired by this book. Along with the idea of surveillance capitalism, and the knowledge of worker exploitation in data curation, and embedded biases into algorithms, you could read this as entertainment, but you could also read it as a plausible warning about the rise of corporate city states and how much culture and technology co-evolve. I recommend that anyone interested in AI read this book, and my recommendation is not to do with the questionable literary value of the novel. :(
	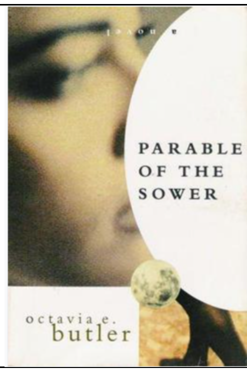	This dystopia isn’t about AI, but it might be the most relevant book on this list. I read an interview with Butler (the author) where she was asked about how she comes up with her ideas and she responded, as a paraphrase, that she simply notices the current societal problem that is being ignored, and then she ages it by 30 years. Interesting approach. The problem she identified was with climate change (AI has a lot of environmental impacts, especially surrounding potable water), inequality, and political unraveling along with subsequent violence. It reminded me that the future is not just about gadgets, but about survival and imagination. I have not yet read the sequel to this novel, but it is on my near-future list!

Do we want to include a list of academic books and papers, too?